

# Problem 1

## Restricted Boltzmann Machine:

We denote the visible features as  $\mathbf{v}$  and the hidden features as  $\mathbf{h}$ . The network of weights between them is  $\mathbf{W}$  with the biases  $\mathbf{c}$  and  $\mathbf{b}$ .

Defining the energy over  $\mathbf{v}$  and  $\mathbf{h}$  as:

$$\begin{aligned} E(\mathbf{v}, \mathbf{h}) &= -\mathbf{h}^T \mathbf{W} \mathbf{v} - \mathbf{c}^T \mathbf{v} - \mathbf{b}^T \mathbf{h} \\ &= -\sum_j \sum_k W_{j,k} h_j v_k - \sum_k c_k v_k - \sum_j b_j h_j \end{aligned}$$

This translates directly to the following free energy formula:  $\mathcal{F}(v) = -c'v - \sum_i \log \sum_{h_i} e^{h_i(b_i + W_i v)}$

Joint probability over  $\mathbf{v}$  and  $\mathbf{h}$  is:

$$\begin{aligned} p(\mathbf{v}, \mathbf{h}) &= \exp(-E(\mathbf{v}, \mathbf{h}))/Z \\ &= \exp(\mathbf{h}^T \mathbf{W} \mathbf{v} + \mathbf{c}^T \mathbf{v} + \mathbf{b}^T \mathbf{h})/Z \\ &= \exp(\mathbf{h}^T \mathbf{W} \mathbf{v}) \exp(\mathbf{c}^T \mathbf{v}) \exp(\mathbf{b}^T \mathbf{h})/Z \\ &= \prod_j \prod_k \exp(W_{j,k} h_j v_k) \prod_k \exp(c_k v_k) \prod_j \exp(b_j h_j) \end{aligned}$$

Because of the specific structure of RBMs, visible and hidden units are conditionally independent. Thus we can say:

$$\begin{aligned} p(h|v) &= \prod_i p(h_i|v) \\ p(v|h) &= \prod_j p(v_j|h) \end{aligned}$$

$$\begin{aligned} \text{RBMs with binary units: } P(h_i = 1|v) &= \sigma(b_i + W_i v) \\ P(v_j = 1|h) &= \sigma(c_j + W'_j h) \end{aligned}$$

Thus the Free energy term reduces to:

$$\mathcal{F}(v) = -c'v - \sum_i \log(1 + e^{(b_i + W_i v)})$$

We aim to maximise the log likelihood:

$$\frac{1}{N} \sum_n \log p(\mathbf{v}^{(n)})$$

$N$  represents no. of training examples.

$$p(\mathbf{v}) = \sum_{\mathbf{h} \in \{0,1\}^H} p(\mathbf{v}, \mathbf{h}) = \sum_{\mathbf{h} \in \{0,1\}^H} \exp(-E(\mathbf{v}, \mathbf{h}))/Z$$

Assuming  $\theta$  be the trainable weights and biases ( $W, c$  and  $b$ ) for stochastic gradient descent we will require:

$$\frac{\partial -\log p(v; \theta)}{\partial \theta} = -\frac{\partial \log \sum_h e^{-E(v, h)}}{\partial \theta} + \frac{\partial \log Z}{\partial \theta}$$

**For the positive phase:**

$$-\frac{\partial \log \sum_h e^{-E(v, h)}}{\partial \theta} = -\frac{1}{\sum_h e^{-E(v, h)}} \sum_h \frac{\partial e^{-E(v, h)}}{\partial \theta} = -\frac{1}{\sum_h e^{-E(v, h)}} \sum_h e^{-E(v, h)} \frac{\partial -E(v, h)}{\partial \theta}$$

After moving the first fraction into the summation, we keep applying the chain rule so to cancel out the negative sign.

$$\sum_h \frac{e^{-E(v, h)}}{\sum_h e^{-E(v, h)}} \frac{\partial E(v, h)}{\partial \theta}$$

Dividing  $Z$  from numerator and denominator of the first fraction simultaneously :

$$\sum_h \frac{\frac{e^{-E(v, h)}}{Z}}{\frac{\sum_h e^{-E(v, h)}}{Z}} \frac{\partial E(v, h)}{\partial \theta} = \sum_h \frac{P(v, h)}{P(v)} \frac{\partial E(v, h)}{\partial \theta} = \sum_h P(h|v) \frac{\partial E(v, h)}{\partial \theta}$$

Thus the final form of the positive phase is:

$$E_{h \sim P(h|v)} \frac{\partial E(v, h)}{\partial \theta}$$

**For the negative phase :**

$$\frac{\partial \log Z}{\partial \theta} = \frac{1}{Z} \frac{\partial Z}{\partial \theta}$$

After substituting  $Z = \sum_{x, h} e^{-E(x, h)}$  and applying chain rule:

$$= \frac{1}{Z} \sum_{h', v'} \frac{\partial e^{-E(h', v')}}{\partial \theta} = \frac{1}{Z} \sum_{h', v'} e^{-E(h', v')} \frac{\partial -E(h', v')}{\partial \theta} = -\frac{1}{Z} \sum_{h', v'} e^{-E(h', v')} \frac{\partial E(h', v')}{\partial \theta}$$

Moving the first fraction into summation so to use joint distribution:

$$= -\sum_{h', v'} \frac{e^{-E(h', v')}}{Z} \frac{\partial E(h', v')}{\partial \theta} = -\sum_{h', v'} P(h', v') \frac{\partial E(h', v')}{\partial \theta}$$

We get the final form of negative phase to be:

$$-E_{h', v' \sim P(h, v)} \frac{\partial E(h', v')}{\partial \theta}$$

We get the following learning rule:

$$\frac{\partial -\log P(v; \theta)}{\partial \theta} = E_{h \sim P(h|v)} \frac{\partial E(v, h)}{\partial \theta} - E_{h', v' \sim P(h, v)} \frac{\partial E(h', v')}{\partial \theta}$$

The parameter updation rule thus becomes:

$$\theta(t) = \theta(t-1) + \eta \frac{\partial \log P(v)}{\partial \theta(t-1)}$$

where  $\eta$  is the learning rate.

Even though the positive phase is tractible, the negative phase is intractible as there are infinite free choices over  $\mathbf{v}$  and  $\mathbf{h}$ . Thus we use **Contrastive Divergence Algorithm** to solve this problem.

Replace estimation over  $(\mathbf{v})$  by a point estimation at  $\tilde{\mathbf{v}} \sim p(\mathbf{v})$  obtained from Gibbs sampling starting at  $\mathbf{v}^{(t)}$

sample  $(\mathbf{h}^1 \sim p(\mathbf{h}|\mathbf{v}^{(t)})$

sample  $(\mathbf{v}^1 \sim p(\mathbf{v}|\mathbf{h}^1)$

This goes on till..

sample  $(\mathbf{h}^k \sim p(\mathbf{h}|\mathbf{v}^{(k-1)})$

sample  $(\mathbf{v} \sim p(\mathbf{v}|\mathbf{h}^{k-1})$

Thus we can write:

$$\begin{aligned} \frac{\partial -\log P(\mathbf{v}; \theta)}{\partial \theta} &= \mathbf{E}_{\mathbf{h}} \left[ \frac{\partial E(\mathbf{v}^{(t)}, \mathbf{h})}{\partial \theta} | \mathbf{v}^{(t)} \right] - \mathbf{E}_{\mathbf{v}, \mathbf{h}} \left[ \frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial \theta} \right] \\ &= \mathbf{E}_{\mathbf{h}} \left[ \frac{\partial E(\mathbf{v}^{(t)}, \mathbf{h})}{\partial \theta} | \mathbf{v}^{(t)} \right] - \mathbf{E}_{\mathbf{h}} \left[ \frac{\partial E(\tilde{\mathbf{v}}, \mathbf{h})}{\partial \theta} \right] \end{aligned}$$

This solves the intractible problem and we can go ahead with parameter updation.

The final update rules are:

$$\mathbf{W} = \mathbf{W} + \alpha(\mathbf{h}(\mathbf{v}^{(t)})\mathbf{v}^{(t)T} - \mathbf{h}(\tilde{\mathbf{v}})\tilde{\mathbf{v}}^T)$$

$$\mathbf{b} = \mathbf{b} + \alpha(\mathbf{h}(\mathbf{v}^{(t)}) - \mathbf{h}(\tilde{\mathbf{v}}))$$

$$\mathbf{c} = \mathbf{c} + \alpha(\mathbf{v}^{(t)} - \tilde{\mathbf{v}})$$

where  $\mathbf{h}(\mathbf{v}) = p(\mathbf{h}|\mathbf{v})$ .

```
In [0]: import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
import os
import keras.datasets
```

```
In [4]: print(tf.__version__)
```

2.2.0-rc3

```
In [0]: import random
import gzip, struct

class DataSet:
    batch_index = 0

    def __init__(self, flag, batch_size = None, one_hot = False, se
```

```

ed = 0):

    if flag=='train':
        (X,Y),(_,_) = keras.datasets.mnist.load_data()
    if flag=='test':
        (_,_), (X,Y) = keras.datasets.mnist.load_data()

    shape = X.shape
    X = X.reshape([shape[0], shape[1] * shape[2]])
    self.X = X.astype(np.float)/255
    self.size = self.X.shape[0]
    if batch_size == None:
        self.batch_size = self.size
    else:
        self.batch_size = batch_size
    # abandon last few samples
    self.batch_num = int(self.size / self.batch_size)
    # shuffle samples
    np.random.seed(seed)
    np.random.shuffle(self.X)
    np.random.seed(seed)
    np.random.shuffle(Y)
    self.one_hot = one_hot
    if one_hot:
        y_vec = np.zeros((len(Y), 10), dtype=np.float)
        for i, label in enumerate(Y):
            y_vec[i, Y[i]] = 1.0
        self.Y = y_vec
    else:
        self.Y = Y

def next_batch(self):
    start = self.batch_index * self.batch_size
    end = (self.batch_index + 1) * self.batch_size
    self.batch_index = (self.batch_index + 1) % self.batch_num
    if self.one_hot:
        return self.X[start:end, :], self.Y[start:end, :]
    else:
        return self.X[start:end, :], self.Y[start:end]

def sample_batch(self):
    index = random.randrange(self.batch_num)
    start = index * self.batch_size
    end = (index + 1) * self.batch_size
    if self.one_hot:
        return self.X[start:end, :], self.Y[start:end, :]
    else:
        return self.X[start:end, :], self.Y[start:end]

```

```
In [6]: mnist = DataSet('test', 2)
print('batch index: %d' % mnist.batch_index)
X, Y = mnist.next_batch()
print('batch index: %d' % mnist.batch_index)
print('X:')
print(X)
print('Y:')
print(Y)
```

```
Downloading data from https://s3.amazonaws.com/img-datasets/mnist.
npz
11493376/11490434 [=====] - 1s 0us/step
batch index: 0
batch index: 1
X:
[[0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]]
Y:
[8 7]
```

```
In [0]: def weight(shape, name='weights'):
        return tf.Variable(tf.random.truncated_normal(shape, stddev=0.1), name=name)

def bias(shape, name='biases'):
    return tf.Variable(tf.constant(0.1, shape=shape), name=name)
```

```
In [0]: class RBM:
        i = 0 # flipping index for computing pseudo likelihood

        def __init__(self, n_hidden, n_visible=784, k=30, momentum=False):
            self.n_visible = n_visible
            self.n_hidden = n_hidden
            self.k = k

            self.lr = tf.compat.v1.placeholder(tf.float32)
            if momentum:
                self.momentum = tf.compat.v1.placeholder(tf.float32)
            else:
                self.momentum = 0.0
            self.w = weight([n_visible, n_hidden], 'w')
            self.hb = bias([n_hidden], 'hb')
            self.vb = bias([n_visible], 'vb')

            self.w_v = tf.Variable(tf.zeros([n_visible, n_hidden]), dtype=tf.float32)
            self.hb_v = tf.Variable(tf.zeros([n_hidden]), dtype=tf.float32)
            self.vb_v = tf.Variable(tf.zeros([n_visible]), dtype=tf.float32)
```

```

def propup(self, visible):
    pre_sigmoid_activation = tf.matmul(visible, self.w) + self.b
    return tf.nn.sigmoid(pre_sigmoid_activation)

def proplean(self, hidden):
    pre_sigmoid_activation = tf.matmul(hidden, tf.transpose(self.w)) + self.b
    return tf.nn.sigmoid(pre_sigmoid_activation)

def sample_h_given_v(self, v_sample):
    h_props = self.propup(v_sample)
    h_sample = tf.nn.relu(tf.sign(h_props - tf.random.uniform(tf.shape(h_props))))
    return h_sample

def sample_v_given_h(self, h_sample):
    v_props = self.proplean(h_sample)
    v_sample = tf.nn.relu(tf.sign(v_props - tf.random.uniform(tf.shape(v_props))))
    return v_sample

def CD_k(self, visibles):
    # k steps gibbs sampling
    v_samples = visibles
    h_samples = self.sample_h_given_v(v_samples)
    for i in range(self.k):
        v_samples = self.sample_v_given_h(h_samples)
        h_samples = self.sample_h_given_v(v_samples)

    h0_props = self.propup(visibles)
    w_positive_grad = tf.matmul(tf.transpose(visibles), h0_props)
    w_negative_grad = tf.matmul(tf.transpose(v_samples), h_samples)
    w_grad = (w_positive_grad - w_negative_grad) / tf.compat.v1.to_float(tf.shape(visibles)[0])
    hb_grad = tf.reduce_mean(h0_props - h_samples, 0)
    vb_grad = tf.reduce_mean(visibles - v_samples, 0)
    return w_grad, hb_grad, vb_grad

def learn(self, visibles):
    w_grad, hb_grad, vb_grad = self.CD_k(visibles)
    # compute new velocities
    new_w_v = self.momentum * self.w_v + self.lr * w_grad
    new_hb_v = self.momentum * self.hb_v + self.lr * hb_grad
    new_vb_v = self.momentum * self.vb_v + self.lr * vb_grad
    # update parameters
    update_w = tf.compat.v1.assign(self.w, self.w + new_w_v)
    update_hb = tf.compat.v1.assign(self.hb, self.hb + new_hb_v)
    update_vb = tf.compat.v1.assign(self.vb, self.vb + new_vb_v)

```

```

        # update velocities
        update_w_v = tf.compat.v1.assign(self.w_v, new_w_v)
        update_hb_v = tf.compat.v1.assign(self.hb_v, new_hb_v)
        update_vb_v = tf.compat.v1.assign(self.vb_v, new_vb_v)

        return [update_w, update_hb, update_vb, update_w_v, update_hb_v, update_vb_v]

    def sampler(self, visibles, steps=200):
        v_samples = visibles
        for step in range(steps):
            v_samples = self.sample_v_given_h(self.sample_h_given_v(v_samples))
        return v_samples

    def free_energy(self, visibles):
        first_term = tf.matmul(visibles, tf.reshape(self.vb, [tf.shape(self.vb)[0], 1]))
        second_term = tf.reduce_sum(tf.math.log(1 + tf.exp(self.hb + tf.matmul(visibles, self.w))), axis=1)
        return - first_term - second_term

    def pseudo_likelihood(self, visibles):
        x = tf.round(visibles)
        x_fe = self.free_energy(x)
        split0, split1, split2 = tf.split(x, [self.i, 1, tf.shape(x)[1] - self.i - 1], 1)
        xi = tf.concat([split0, 1 - split1, split2], 1)
        self.i = (self.i + 1) % self.n_visible
        xi_fe = self.free_energy(xi)
        return tf.reduce_mean(self.n_visible * tf.math.log(tf.nn.sigmoid(xi_fe - x_fe)), axis=0)

```

```
In [0]: def show_images(images, size):

    img = (images + 1.0) / 2.0
    h, w = img.shape[1], img.shape[2]

    merge_img = np.zeros((h * size[0], w * size[1]))

    for idx, image in enumerate(images):
        i = idx % size[1]
        j = idx // size[1]
        merge_img[j*h:j*h+h, i*w:i*w+w] = image

    merge_img = (merge_img * 255 / np.max(merge_img)).astype('uint8')
    plt.imshow(merge_img, cmap='gray')
    plt.show()

    return
```

```
In [0]: def train(train_data, hidden_nodes, epoches):

    x = tf.compat.v1.placeholder(tf.float32, shape=[None, 784])
    noise_x, _ = train_data.sample_batch()
    # noise_x = tf.random_normal([train_data.batch_size, 784])
    rbm = RBM(hidden_nodes)
    step = rbm.learn(x)
    sampler = rbm.sampler(x)

    saver = tf.compat.v1.train.Saver()

    with tf.compat.v1.Session() as sess:
        init = tf.compat.v1.global_variables_initializer()
        sess.run(init)
        #mean_cost = []
        epoch = 1
        for i in range(epoches * train_data.batch_num):
            # draw samples
            if i % train_data.batch_num == 0:
                samples = sess.run(sampler, feed_dict = {x: noise_x})
                samples = samples.reshape([train_data.batch_size, 28, 28])
                show_images(samples, [8, 8])

            batch_x, _ = train_data.next_batch()
            sess.run(step, feed_dict = {x: batch_x, rbm.lr: 0.1})

            if i is not 0 and train_data.batch_index is 0:
```



```
        print('Epoch ', epoch)
        epoch += 1

    print('Testing !')

    x_test, _ = test_data.sample_batch()

    x_test1 = np.copy(x_test)
    x_test2 = np.copy(x_test)
    x_test3 = np.copy(x_test)

    l1, l2, l3 = [], [], []
    for j in range(len(x_test)):
        i, = np.nonzero(x_test[j])
        ix1 = np.random.choice(i, int(np.floor(0.2 * len(i))), replace=False)
        ix2 = np.random.choice(i, int(np.floor(0.5 * len(i))), replace=False)
        ix3 = np.random.choice(i, int(np.floor(0.8 * len(i))), replace=False)

        l1.append(ix1)
        l2.append(ix2)
        l3.append(ix3)

    print('After removing 20% pixels')

    j = 0
    for ix in l1:
        for loc in ix:
            x_test1[j][loc] = 0.0
        j += 1

    samples = sess.run(sampler, feed_dict = {x: x_test1})
    samples = samples.reshape([train_data.batch_size, 28, 28])
    show_images(samples, [8, 8])

    print('After removing 50% pixels')

    j = 0
    for ix in l2:
        for loc in ix:
            x_test2[j][loc] = 0.0
        j += 1

    samples = sess.run(sampler, feed_dict = {x: x_test2})
    samples = samples.reshape([train_data.batch_size, 28, 28])
    show_images(samples, [8, 8])
```

```

print('After removing 80% pixels')

j=0
for ix in l3:
    for loc in ix:
        x_test3[j][loc]=0.0
        j+=1

samples = sess.run(sampler, feed_dict = {x: x_test3})
samples = samples.reshape([train_data.batch_size, 28, 28])
show_images(samples, [8, 8])

```

```

In [0]: train_data = DataSet('train', batch_size=64, one_hot=True)
        test_data = DataSet('test', batch_size=64, one_hot=True)

```

**20 hidden nodes training.**

**Followed by testing: removing 20%, 50% and 80% pixels.**

```

In [20]: tf.compat.v1.disable_eager_execution()
        train(train_data, 20,30)

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/resource\_variable\_ops.py:1666: calling BaseResourceVariable.\_\_init\_\_ (from tensorflow.python.ops.resource\_variable\_ops) with constraint is deprecated and will be removed in a future version.


Instructions for updating:

If using Keras pass `*_constraint` arguments to layers.


WARNING:tensorflow:From <ipython-input-8-9bfcc859e862>:51: `to_float` (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:


Use ``tf.cast`` instead.




Epoch 1




Epoch 2




Epoch 3




Epoch 4




Epoch 5




Epoch 6




Epoch 7




Epoch 8




Epoch 9




Epoch 10




Epoch 11




Epoch 12




Epoch 13




Epoch 14




Epoch 15




Epoch 16



Epoch 17




Epoch 18




Epoch 19







Epoch 20




Epoch 21




Epoch 22




Epoch 23




Epoch 24




Epoch 25




Epoch 26




Epoch 27



Epoch 28




Epoch 29




Epoch 30

Testing !


After removing 20% pixels



After removing 50% pixels




After removing 80% pixels




**100 hidden nodes training.**

**Followed by testing: removing 20%, 50% and 80% pixels.**


```
In [21]: tf.compat.v1.disable_eager_execution()
          train(train_data, 100,30)
```




Epoch 1




Epoch 2




Epoch 3




Epoch 4




Epoch 5




Epoch 6



Epoch 7




Epoch 8




Epoch 9







Epoch 10




Epoch 11




Epoch 12




Epoch 13




Epoch 14




Epoch 15




Epoch 16




Epoch 17




Epoch 18




Epoch 19




Epoch 20




Epoch 21




Epoch 22




Epoch 23




Epoch 24




Epoch 25




Epoch 26




Epoch 27



Epoch 28




Epoch 29




Epoch 30

Testing !


After removing 20% pixels



After removing 50% pixels



After removing 80% pixels







**500 hidden nodes training.**

**Followed by testing: removing 20%, 50% and 80% pixels.**


```
In [22]: tf.compat.v1.disable_eager_execution()  
train(train_data, 500,30)
```




Epoch 1




Epoch 2




Epoch 3




Epoch 4




Epoch 5




Epoch 6




Epoch 7




Epoch 8




Epoch 9




Epoch 10




Epoch 11




Epoch 12




Epoch 13




Epoch 14




Epoch 15




Epoch 16




Epoch 17




Epoch 18




Epoch 19




Epoch 20



Epoch 21




Epoch 22




Epoch 23







Epoch 24




Epoch 25




Epoch 26




Epoch 27



Epoch 28




Epoch 29




Epoch 30

Testing !


After removing 20% pixels



After removing 50% pixels



After removing 80% pixels



## Problem 2

### VAE derivation:

Let,

$X$ : data that we want to model

$z$ : latent variable

$P(X)$ : probability distribution of the data

$P(z)$ : probability distribution of latent variable

$P(X|z)$ : distribution of generating data given latent variable

Our objective here is to model the data, hence we want to find  $P(X)$ . Using the law of probability, we could find it in relation with  $z$  as follows:

$$P(X) = \int P(X|z)P(z)dz$$

But the problem arises since the integral is intractable. So our aim is to approximate  $P(z|X)$  by  $Q(z|X)$  using variational inference.

let's say we want to infer  $P(z|X)$  using  $Q(z|X)$ . The KL divergence then formulated as follows:

$$D_{KL} Q(z|X) || P(z|X) = \sum_z Q(z|X) \log \frac{Q(z|X)}{P(z|X)}$$

Using Bayes Rule:

$$= E[\log \frac{Q(z|X)}{P(z|X)}]$$

$$= E[\log Q(z|X) - \log P(z|X)]$$

$$D_{KL} Q(z|X) || P(z|X) = E[\log Q(z|X) - \log \frac{P(X|z)P(z)}{P(X)}]$$

$$= E[\log Q(z|X) - \log P(X|z) - \log P(z) + \log P(X)]$$

As  $P(X)$  is independent of  $z$  we move it outside the expectation term, this gives us the following:

$$D_{KL}[Q(z|X)||P(z|X)] = E[\log Q(z|X) - \log P(X|z) - \log P(z)] + \log P(X)$$

$$D_{KL}[Q(z|X)||P(z|X)] - \log P(X) = E[\log Q(z|X) - \log P(X|z) - \log P(z)]$$

On looking at the right hand side of the equation, we would notice that it could be rewritten as another KL divergence:

$$\begin{aligned} \log P(X) - D_{KL}[Q(z|X)||P(z|X)] &= E[\log P(X|z) - (\log Q(z|X) - \log P(z))] \\ &= E[\log P(X|z)] - E[\log Q(z|X) - \log P(z)] \\ &= E[\log P(X|z)] - D_{KL}[Q(z|X)||P(z)] \end{aligned}$$

Thus we get the **VAE objective function** as:

$$\log P(X) - D_{KL}[Q(z|X)||P(z|X)] = E[\log P(X|z)] - D_{KL}[Q(z|X)||P(z)]$$

We have the following 3 steps:

1. Learn  $Q(z|X)$  by transforming data  $X$  into  $z$  i.e. the Encoder module.
2. Learn  $P(z)$  i.e. distribution of latent space.
3. Generate new data from  $z$   $P(X|z)$  i.e. the Decoder module.

Maximizing  $E[\log P(X|z)]$  is a maximum likelihood estimation. However for the other term  $D_{KL}[Q(z|X)||P(z)]$ ,  $P(z)$  is the latent variable distribution. When sampling  $P(z)$  the easiest choice is  $N(0,1)$ . Hence,  $Q(z|X)$  must be close to  $N(0,1)$ . Also then, the KL divergence between those two distribution could be computed in closed form.

$$\begin{aligned} D_{KL}[N(\mu(X), \Sigma(X))||N(0, 1)] &= \frac{1}{2}(\sum_k \Sigma(X) + \sum_k \mu^2(X) - \sum_k 1 - \log \prod_k \Sigma(X)) \\ &= \frac{1}{2} \sum_k (\Sigma(X) + \mu^2(X) - 1 - \log \Sigma(X)) \end{aligned}$$

In practice, however, it's better to model  $\sigma(X)$  as  $\log \sigma(X)$ , as it is more numerically stable to take exponent compared to computing log. Hence, our final KL divergence term is:

$$D_{KL}[N(\mu(X), \Sigma(X))||N(0, 1)] = \frac{1}{2} \sum_k (\exp(\Sigma(X)) + \mu^2(X) - 1 - \Sigma(X))$$

**Reparameterization trick** is used to divert the non-differentiable operation out of the network, so that, even though we still involve a thing that is non-differentiable, at least it is out of the network, hence the network could still be trained. The reparameterization trick is as follows. If we sample from a standard normal distribution, we could convert it to any Gaussian we want if we know the mean and the variance. Hence we could implement our sampling operation of  $z$  by:

$$z = \mu(X) + \Sigma^{\frac{1}{2}}(X)\epsilon$$

where  $\epsilon \sim N(0, 1)$

During backpropagation, we don't care anymore with the sampling process, as it is now outside of the network, i.e. doesn't depend on anything in the net, hence the gradient won't flow through it.

```
In [0]: from keras.layers import Lambda, Input, Dense
        from keras.models import Model
        from keras.datasets import mnist
        from keras.losses import mse, binary_crossentropy
        from keras.utils import plot_model
        from keras import backend as K

        import numpy as np
        import matplotlib.pyplot as plt
        import argparse
        import os
```

```
In [0]: def sampling(args):

        z_mean, z_log_var = args
        batch = K.shape(z_mean)[0]
        dim = K.int_shape(z_mean)[1]
        # by default, random_normal has mean = 0 and std = 1.0
        epsilon = K.random_normal(shape=(batch, dim))
        return z_mean + K.exp(0.5 * z_log_var) * epsilon
```

```
In [0]: def plot_results(models,
                        data,
                        batch_size=128,
                        model_name="vae_mnist"):

    encoder, decoder = models
    x_test, y_test = data

    # display a 30x30 2D manifold of digits
    n = 30
    digit_size = 28
    figure = np.zeros((digit_size * n, digit_size * n))
    # linearly spaced coordinates corresponding to the 2D plot
    # of digit classes in the latent space
    grid_x = np.linspace(-3, 3, n)
    grid_y = np.linspace(-3, 3, n)[::-1]

    for i, yi in enumerate(grid_y):
        for j, xi in enumerate(grid_x):
            z_sample = np.array([[xi, yi]])
            x_decoded = decoder.predict(z_sample)
            digit = x_decoded[0].reshape(digit_size, digit_size)
            figure[i * digit_size: (i + 1) * digit_size,
                  j * digit_size: (j + 1) * digit_size] = digit

    plt.figure(figsize=(10, 10))
    start_range = digit_size // 2
    end_range = (n - 1) * digit_size + start_range + 1
    pixel_range = np.arange(start_range, end_range, digit_size)
    sample_range_x = np.round(grid_x, 1)
    sample_range_y = np.round(grid_y, 1)
    plt.xticks(pixel_range, sample_range_x)
    plt.yticks(pixel_range, sample_range_y)
    plt.imshow(figure, cmap='Greys_r')
    plt.show()
```

```
In [0]: def vae(latent_dim):

    inputs = Input(shape=input_shape, name='encoder_input')
    x = Dense(intermediate_dim, activation='relu')(inputs)
    z_mean = Dense(latent_dim, name='z_mean')(x)
    z_log_var = Dense(latent_dim, name='z_log_var')(x)

    z = Lambda(sampling, output_shape=(latent_dim,), name='z')([z_mean, z_log_var])

    # instantiate encoder model
    encoder = Model(inputs, [z_mean, z_log_var, z], name='encoder')
    encoder.summary()
```

```
# build decoder model
latent_inputs = Input(shape=(latent_dim,), name='z_sampling')
x = Dense(intermediate_dim, activation='relu')(latent_inputs)
outputs = Dense(original_dim, activation='sigmoid')(x)

# instantiate decoder model
decoder = Model(latent_inputs, outputs, name='decoder')
decoder.summary()

# instantiate VAE model
outputs = decoder(encoder(inputs)[2])
vae = Model(inputs, outputs, name='vae_mlp')
vae.summary()

models = (encoder, decoder)
data = (x_test, y_test)

reconstruction_loss = mse(inputs, outputs)
reconstruction_loss *= original_dim
kl_loss = 1 + z_log_var - K.square(z_mean) - K.exp(z_log_var)
kl_loss = K.sum(kl_loss, axis=-1)
kl_loss *= -0.5
vae_loss = K.mean(reconstruction_loss + kl_loss)
vae.add_loss(vae_loss)
vae.compile(optimizer='adam')
vae.summary()

vae.fit(x_train, epochs=epochs, batch_size=batch_size, validation_data=(x_test, None))

if latent_dim==2:
    plot_results(models, data, batch_size=batch_size, model_name="vae_mlp")

return vae
```



```
In [0]: # MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()

image_size = x_train.shape[1]
original_dim = image_size * image_size
x_train = np.reshape(x_train, [-1, original_dim])
x_test = np.reshape(x_test, [-1, original_dim])
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255

# network parameters
input_shape = (original_dim, )
intermediate_dim = 256
batch_size = 128
latent_dim = 2
epochs = 30
```

```
In [0]: original_dim, intermediate_dim
```

```
Out[0]: (784, 256)
```

### Code units=2 and generated images by varying latent variables from -3 to 3

```
In [0]: model=vae(2)
```

Model: "encoder"

Layer (type) connected to	Output Shape	Param #	C
=====			
encoder_input (InputLayer)	(None, 784)	0	
dense_1 (Dense) ncoder_input[0][0]	(None, 256)	200960	e
z_mean (Dense) ense_1[0][0]	(None, 2)	514	d
z_log_var (Dense) ense_1[0][0]	(None, 2)	514	d
z (Lambda) _mean[0][0]	(None, 2)	0	z
_log_var[0][0]			z

```

=====
Total params: 201,988
Trainable params: 201,988
Non-trainable params: 0

```

---

Model: "decoder"

Layer (type)	Output Shape	Param #
z_sampling (InputLayer)	(None, 2)	0
dense_2 (Dense)	(None, 256)	768
dense_3 (Dense)	(None, 784)	201488

```

=====
Total params: 202,256
Trainable params: 202,256
Non-trainable params: 0

```

---

Model: "vae\_mlp"

Layer (type)	Output Shape	Param #
encoder_input (InputLayer)	(None, 784)	0
encoder (Model)	[(None, 2), (None, 2), (N	201988
decoder (Model)	(None, 784)	202256

```

=====
Total params: 404,244
Trainable params: 404,244
Non-trainable params: 0

```

---

Model: "vae\_mlp"

Layer (type)	Output Shape	Param #
encoder_input (InputLayer)	(None, 784)	0
encoder (Model)	[(None, 2), (None, 2), (N	201988
decoder (Model)	(None, 784)	202256

```

=====
Total params: 404,244
Trainable params: 404,244
Non-trainable params: 0

```

---

```
/usr/local/lib/python3.6/dist-packages/keras/engine/training_utils.py:819: UserWarning: Output decoder missing from loss dictionary. We assume this was done on purpose. The fit and evaluate APIs will not be expecting any data to be passed to decoder.
```

```
'be expecting any data to be passed to {0}.'.format(name))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

```
60000/60000 [=====] - 3s 58us/step - loss : 55.0477 - val_loss: 45.6592
```

Epoch 2/30

```
60000/60000 [=====] - 2s 33us/step - loss : 44.0658 - val_loss: 43.2927
```

Epoch 3/30

```
60000/60000 [=====] - 2s 33us/step - loss : 42.7544 - val_loss: 42.5227
```

Epoch 4/30

```
60000/60000 [=====] - 2s 33us/step - loss : 42.1759 - val_loss: 41.9824
```

Epoch 5/30

```
60000/60000 [=====] - 2s 33us/step - loss : 41.7146 - val_loss: 41.5494
```

Epoch 6/30

```
60000/60000 [=====] - 2s 33us/step - loss : 41.3244 - val_loss: 41.1091
```

Epoch 7/30

```
60000/60000 [=====] - 2s 33us/step - loss : 40.9802 - val_loss: 40.8690
```

Epoch 8/30

```
60000/60000 [=====] - 2s 34us/step - loss : 40.6434 - val_loss: 40.5432
```

Epoch 9/30

```
60000/60000 [=====] - 2s 33us/step - loss : 40.3415 - val_loss: 40.3838
```

Epoch 10/30

```
60000/60000 [=====] - 2s 33us/step - loss : 40.0907 - val_loss: 40.0736
```

Epoch 11/30

```
60000/60000 [=====] - 2s 33us/step - loss : 39.8978 - val_loss: 39.9238
```

Epoch 12/30

```
60000/60000 [=====] - 2s 33us/step - loss : 39.6876 - val_loss: 39.7079
```

Epoch 13/30

```
60000/60000 [=====] - 2s 33us/step - loss : 39.5086 - val_loss: 39.6411
```

Epoch 14/30


```
60000/60000 [=====] - 2s 33us/step - loss : 39.3637 - val_loss: 39.4656
```

Epoch 15/30

```
60000/60000 [=====] - 2s 33us/step - loss : 39.2222 - val_loss: 39.2878
```

Epoch 16/30

```
60000/60000 [=====] - 2s 33us/step - loss
: 39.0595 - val_loss: 39.2510
Epoch 17/30
60000/60000 [=====] - 2s 33us/step - loss
: 38.9537 - val_loss: 39.1338
Epoch 18/30
60000/60000 [=====] - 2s 32us/step - loss
: 38.8503 - val_loss: 39.0146
Epoch 19/30
60000/60000 [=====] - 2s 33us/step - loss
: 38.7450 - val_loss: 38.9454
Epoch 20/30
60000/60000 [=====] - 2s 33us/step - loss
: 38.6451 - val_loss: 38.8135
Epoch 21/30
60000/60000 [=====] - 2s 32us/step - loss
: 38.5462 - val_loss: 38.7438
Epoch 22/30
60000/60000 [=====] - 2s 33us/step - loss
: 38.4631 - val_loss: 38.7279
Epoch 23/30
60000/60000 [=====] - 2s 33us/step - loss
: 38.3945 - val_loss: 38.5867
Epoch 24/30
60000/60000 [=====] - 2s 32us/step - loss
: 38.3223 - val_loss: 38.6438
Epoch 25/30
60000/60000 [=====] - 2s 33us/step - loss
: 38.2469 - val_loss: 38.4917
Epoch 26/30
60000/60000 [=====] - 2s 33us/step - loss
: 38.1658 - val_loss: 38.5023
Epoch 27/30
60000/60000 [=====] - 2s 33us/step - loss
: 38.1239 - val_loss: 38.4634
Epoch 28/30
60000/60000 [=====] - 2s 32us/step - loss
: 38.0551 - val_loss: 38.4107
Epoch 29/30
60000/60000 [=====] - 2s 33us/step - loss
: 38.0106 - val_loss: 38.5096
Epoch 30/30
60000/60000 [=====] - 2s 33us/step - loss
: 37.9651 - val_loss: 38.3257
```




```
In [0]: inference=model.predict(x_test)
```

```
In [0]: inference.shape
```

```
Out[0]: (10000, 784)
```

```
In [0]: num_row = 2
num_col = 5
# plot images
fig, axes = plt.subplots(num_row, num_col, figsize=(1.5*num_col,2*num_row))
for i in range(10):
    ax = axes[i//num_col, i%num_col]
    ax.imshow(inference[i].reshape(28,28), cmap='gray')
plt.tight_layout()
plt.show()
```



### Code units=8

```
In [0]: model=vae(8)
```

Model: "encoder"

Layer (type) connected to	Output Shape	Param #	C
encoder_input (InputLayer)	(None, 784)	0	
dense_4 (Dense) encoder_input[0][0]	(None, 256)	200960	e
z_mean (Dense) dense_4[0][0]	(None, 8)	2056	d
z_log_var (Dense) dense_4[0][0]	(None, 8)	2056	d

```

z (Lambda) (None, 8) 0 z
_mean[0][0]
z
_log_var[0][0]
=====
=====

```

```

Total params: 205,072
Trainable params: 205,072
Non-trainable params: 0

```

```

Model: "decoder"

```

Layer (type)	Output Shape	Param #
z_sampling (InputLayer)	(None, 8)	0
dense_5 (Dense)	(None, 256)	2304
dense_6 (Dense)	(None, 784)	201488

```

Total params: 203,792
Trainable params: 203,792
Non-trainable params: 0

```

```

Model: "vae_mlp"

```

Layer (type)	Output Shape	Param #
encoder_input (InputLayer)	(None, 784)	0
encoder (Model)	[(None, 8), (None, 8), (N 205072	
decoder (Model)	(None, 784)	203792

```

Total params: 408,864
Trainable params: 408,864
Non-trainable params: 0

```

```

Model: "vae_mlp"

```

Layer (type)	Output Shape	Param #
encoder_input (InputLayer)	(None, 784)	0
encoder (Model)	[(None, 8), (None, 8), (N 205072	
decoder (Model)	(None, 784)	203792

```

Total params: 408,864
Trainable params: 408,864
Non-trainable params: 0

```

---

```
/usr/local/lib/python3.6/dist-packages/keras/engine/training_utils.py:819: UserWarning: Output decoder missing from loss dictionary. We assume this was done on purpose. The fit and evaluate APIs will not be expecting any data to be passed to decoder.
```

```
'be expecting any data to be passed to {0}.'.format(name))
```

```
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/30
```

```
60000/60000 [=====] - 4s 61us/step - loss : 48.4368 - val_loss: 37.5004
```

```
Epoch 2/30
```

```
60000/60000 [=====] - 2s 33us/step - loss : 36.1297 - val_loss: 34.7652
```

```
Epoch 3/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 34.4447 - val_loss: 33.6796
```

```
Epoch 4/30
```

```
60000/60000 [=====] - 2s 33us/step - loss : 33.5777 - val_loss: 33.0398
```

```
Epoch 5/30
```

```
60000/60000 [=====] - 2s 33us/step - loss : 32.9973 - val_loss: 32.4594
```

```
Epoch 6/30
```

```
60000/60000 [=====] - 2s 33us/step - loss : 32.5824 - val_loss: 32.2101
```

```
Epoch 7/30
```

```
60000/60000 [=====] - 2s 33us/step - loss : 32.2896 - val_loss: 32.0000
```

```
Epoch 8/30
```

```
60000/60000 [=====] - 2s 33us/step - loss : 31.9871 - val_loss: 31.7186
```

```
Epoch 9/30
```

```
60000/60000 [=====] - 2s 33us/step - loss : 31.7781 - val_loss: 31.4970
```

```
Epoch 10/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 31.5954 - val_loss: 31.2821
```

```
Epoch 11/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 31.4011 - val_loss: 31.1281
```

```
Epoch 12/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 31.2556 - val_loss: 31.1085
```

```
Epoch 13/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 31.1052 - val_loss: 30.8670
```

```
Epoch 14/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 30.9381 - val_loss: 30.8271
```

```
Epoch 15/30
```


```
60000/60000 [=====] - 2s 32us/step - loss : 30.8391 - val_loss: 30.5673
```



```
Epoch 16/30
60000/60000 [=====] - 2s 32us/step - loss
: 30.7304 - val_loss: 30.5606
Epoch 17/30
60000/60000 [=====] - 2s 32us/step - loss
: 30.6023 - val_loss: 30.4539
Epoch 18/30
60000/60000 [=====] - 2s 32us/step - loss
: 30.5075 - val_loss: 30.2782
Epoch 19/30
60000/60000 [=====] - 2s 32us/step - loss
: 30.4349 - val_loss: 30.2428
Epoch 20/30
60000/60000 [=====] - 2s 32us/step - loss
: 30.3516 - val_loss: 30.1365
Epoch 21/30
60000/60000 [=====] - 2s 33us/step - loss
: 30.2616 - val_loss: 30.1102
Epoch 22/30
60000/60000 [=====] - 2s 32us/step - loss
: 30.1911 - val_loss: 30.0771
Epoch 23/30
60000/60000 [=====] - 2s 32us/step - loss
: 30.1357 - val_loss: 29.9404
Epoch 24/30
60000/60000 [=====] - 2s 32us/step - loss
: 30.0499 - val_loss: 29.9513
Epoch 25/30
60000/60000 [=====] - 2s 32us/step - loss
: 29.9813 - val_loss: 29.8946
Epoch 26/30
60000/60000 [=====] - 2s 32us/step - loss
: 29.9199 - val_loss: 29.8648
Epoch 27/30
60000/60000 [=====] - 2s 32us/step - loss
: 29.8733 - val_loss: 29.6888
Epoch 28/30
60000/60000 [=====] - 2s 32us/step - loss
: 29.8457 - val_loss: 29.7381
Epoch 29/30
60000/60000 [=====] - 2s 32us/step - loss
: 29.7829 - val_loss: 29.6970
Epoch 30/30
60000/60000 [=====] - 2s 32us/step - loss
: 29.7382 - val_loss: 29.6226
```

```
In [0]: inference=model.predict(x_test)
```

```
In [0]: num_row = 2
num_col = 5
# plot images
fig, axes = plt.subplots(num_row, num_col, figsize=(1.5*num_col,2*num_row))
for i in range(10):
    ax = axes[i//num_col, i%num_col]
    ax.imshow(inference[i].reshape(28,28), cmap='gray')
plt.tight_layout()
plt.show()
```



### Code units=16

```
In [0]: model=vae(16)
```

Model: "encoder"

Layer (type) connected to	Output Shape	Param #	C
encoder_input (InputLayer)	(None, 784)	0	
dense_7 (Dense) encoder_input[0][0]	(None, 256)	200960	e
z_mean (Dense) dense_7[0][0]	(None, 16)	4112	d
z_log_var (Dense) dense_7[0][0]	(None, 16)	4112	d

```

z (Lambda) (None, 16) 0 z
_mean[0][0]
z
_log_var[0][0]
=====
=====

```

```

Total params: 209,184
Trainable params: 209,184
Non-trainable params: 0

```

```

Model: "decoder"

```

Layer (type)	Output Shape	Param #
z_sampling (InputLayer)	(None, 16)	0
dense_8 (Dense)	(None, 256)	4352
dense_9 (Dense)	(None, 784)	201488

```

Total params: 205,840
Trainable params: 205,840
Non-trainable params: 0

```

```

Model: "vae_mlp"

```

Layer (type)	Output Shape	Param #
encoder_input (InputLayer)	(None, 784)	0
encoder (Model)	[(None, 16), (None, 16),	209184
decoder (Model)	(None, 784)	205840

```

Total params: 415,024
Trainable params: 415,024
Non-trainable params: 0

```

```

Model: "vae_mlp"

```

Layer (type)	Output Shape	Param #
encoder_input (InputLayer)	(None, 784)	0
encoder (Model)	[(None, 16), (None, 16),	209184
decoder (Model)	(None, 784)	205840

```

Total params: 415,024
Trainable params: 415,024
Non-trainable params: 0

```

---

```
/usr/local/lib/python3.6/dist-packages/keras/engine/training_utils.py:819: UserWarning: Output decoder missing from loss dictionary. We assume this was done on purpose. The fit and evaluate APIs will not be expecting any data to be passed to decoder.
```

```
'be expecting any data to be passed to {0}.'.format(name))
```

```
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/30
```

```
60000/60000 [=====] - 4s 62us/step - loss : 49.0297 - val_loss: 37.7708
```

```
Epoch 2/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 36.0786 - val_loss: 34.4709
```

```
Epoch 3/30
```

```
60000/60000 [=====] - 2s 33us/step - loss : 33.9707 - val_loss: 33.1215
```

```
Epoch 4/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 32.9443 - val_loss: 32.2671
```

```
Epoch 5/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 32.3295 - val_loss: 31.7344
```

```
Epoch 6/30
```

```
60000/60000 [=====] - 2s 33us/step - loss : 31.9197 - val_loss: 31.4109
```

```
Epoch 7/30
```

```
60000/60000 [=====] - 2s 33us/step - loss : 31.6239 - val_loss: 31.1337
```

```
Epoch 8/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 31.3720 - val_loss: 31.0695
```

```
Epoch 9/30
```

```
60000/60000 [=====] - 2s 31us/step - loss : 31.1935 - val_loss: 30.9267
```

```
Epoch 10/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 31.0544 - val_loss: 30.7454
```

```
Epoch 11/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 30.9226 - val_loss: 30.6007
```

```
Epoch 12/30
```

```
60000/60000 [=====] - 2s 31us/step - loss : 30.7965 - val_loss: 30.4270
```

```
Epoch 13/30
```

```
60000/60000 [=====] - 2s 32us/step - loss : 30.6800 - val_loss: 30.3863
```

```
Epoch 14/30
```

```
60000/60000 [=====] - 2s 31us/step - loss : 30.5851 - val_loss: 30.3613
```


```
Epoch 15/30
```

```
60000/60000 [=====] - 2s 31us/step - loss : 30.5006 - val_loss: 30.2732
```

```
Epoch 16/30
60000/60000 [=====] - 2s 31us/step - loss
: 30.4565 - val_loss: 30.1469
Epoch 17/30
60000/60000 [=====] - 2s 31us/step - loss
: 30.3848 - val_loss: 30.0676
Epoch 18/30
60000/60000 [=====] - 2s 32us/step - loss
: 30.2877 - val_loss: 30.1163
Epoch 19/30
60000/60000 [=====] - 2s 31us/step - loss
: 30.2353 - val_loss: 30.0230
Epoch 20/30
60000/60000 [=====] - 2s 31us/step - loss
: 30.1928 - val_loss: 29.9477
Epoch 21/30
60000/60000 [=====] - 2s 32us/step - loss
: 30.1131 - val_loss: 30.0213
Epoch 22/30
60000/60000 [=====] - 2s 32us/step - loss
: 30.1035 - val_loss: 29.8594
Epoch 23/30
60000/60000 [=====] - 2s 32us/step - loss
: 30.0254 - val_loss: 29.8712
Epoch 24/30
60000/60000 [=====] - 2s 32us/step - loss
: 29.9937 - val_loss: 29.8370
Epoch 25/30
60000/60000 [=====] - 2s 32us/step - loss
: 29.9424 - val_loss: 29.7841
Epoch 26/30
60000/60000 [=====] - 2s 31us/step - loss
: 29.9131 - val_loss: 29.7074
Epoch 27/30
60000/60000 [=====] - 2s 32us/step - loss
: 29.8772 - val_loss: 29.6718
Epoch 28/30
60000/60000 [=====] - 2s 32us/step - loss
: 29.8316 - val_loss: 29.6844
Epoch 29/30
60000/60000 [=====] - 2s 33us/step - loss
: 29.7857 - val_loss: 29.6533
Epoch 30/30
60000/60000 [=====] - 2s 32us/step - loss
: 29.7461 - val_loss: 29.5284
```

```
In [0]: inference=model.predict(x_test)
```

```
In [0]: num_row = 2
num_col = 5
# plot images
fig, axes = plt.subplots(num_row, num_col, figsize=(1.5*num_col,2*num_row))
for i in range(10):
    ax = axes[i//num_col, i%num_col]
    ax.imshow(inference[i].reshape(28,28), cmap='gray')
plt.tight_layout()
plt.show()
```



## Problem 3

```
In [0]: from keras.layers import Dense, Input
from keras.layers import Conv2D, Flatten, Lambda
from keras.layers import Reshape, Conv2DTranspose
from keras.models import Model
from keras.datasets import mnist
from keras.losses import mse, binary_crossentropy
from keras import backend as K

import numpy as np
import matplotlib.pyplot as plt
import argparse
import os
```

Using TensorFlow backend.

### Load MNIST data

```
In [0]: (x_train, y_train), (x_test, y_test) = mnist.load_data()

image_size = x_train.shape[1]

train_len = 10000
test_len = 1000
x_train = np.reshape(x_train[:train_len], [-1, image_size, image_size, 1])
x_test = np.reshape(x_test[:test_len], [-1, image_size, image_size, 1])
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255
y_test = y_test[:test_len]
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 1s 0us/step

```
In [0]: def sampling(args):
        z_mean, z_log_var = args
        batch = K.shape(z_mean)[0]
        dim = K.int_shape(z_mean)[1]
        # by default, random_normal has mean=0 and std=1.0
        epsilon = K.random_normal(shape=(batch, dim))
        return z_mean + K.exp(0.5 * z_log_var) * epsilon
```

```
In [0]: def plot_results(encoder, decoder, x_test, y_test, batch_size=128):

    # display a 30x30 2D manifold of digits
    n = 30
    digit_size = 28
    figure = np.zeros((digit_size * n, digit_size * n))
    # linearly spaced coordinates corresponding to the 2D plot
    # of digit classes in the latent space
    grid_x = np.linspace(-4, 4, n)
    grid_y = np.linspace(-4, 4, n)[::-1]

    for i, yi in enumerate(grid_y):
        for j, xi in enumerate(grid_x):
            z_sample = np.array([[xi, yi]])
            x_decoded = decoder.predict(z_sample)
            digit = x_decoded[0].reshape(digit_size, digit_size)
            figure[i * digit_size: (i + 1) * digit_size,
                  j * digit_size: (j + 1) * digit_size] = digit

    plt.figure(figsize=(10, 10))
    start_range = digit_size // 2
    end_range = n * digit_size + start_range + 1
    pixel_range = np.arange(start_range, end_range, digit_size)
    sample_range_x = np.round(grid_x, 1)
    sample_range_y = np.round(grid_y, 1)
    plt.xticks(pixel_range, sample_range_x)
    plt.yticks(pixel_range, sample_range_y)
    plt.xlabel("z[0]")
    plt.ylabel("z[1]")
    plt.imshow(figure, cmap='Greys_r')
    plt.show()
```

## Hyperparameters

```
In [0]: input_shape = (image_size, image_size, 1)
        batch_size = 32
        kernel_size = 3
        filters = 16
        latent_dim = 2
        epochs = 30
        use_mse = True
```

## Encoder with CNN



```
In [0]: inputs = Input(shape=input_shape, name='encoder_input')
x = inputs
for i in range(2):
    filters *= 2
    x = Conv2D(filters=filters,
               kernel_size=kernel_size,
               activation='relu',
               strides=2,
               padding='same')(x)

# shape info needed to build decoder model
shape = K.int_shape(x)

# generate latent vector  $Q(z/X)$ 
x = Flatten()(x)
x = Dense(16, activation='relu')(x)
z_mean = Dense(latent_dim, name='z_mean')(x)
z_log_var = Dense(latent_dim, name='z_log_var')(x)

# use reparameterization trick to push the sampling out as input
# note that "output_shape" isn't necessary with the TensorFlow back
end
z = Lambda(sampling, output_shape=(latent_dim,), name='z')([z_mean,
z_log_var])

# instantiate encoder model
encoder = Model(inputs, [z_mean, z_log_var, z], name='encoder')
encoder.summary()
```

Model: "encoder"

Layer (type) connected to	Output Shape	Param #	C
=====			
encoder_input (InputLayer)	(None, 28, 28, 1)	0	
=====			
conv2d_1 (Conv2D) encoder_input[0][0]	(None, 14, 14, 32)	320	e
=====			
conv2d_2 (Conv2D) conv2d_1[0][0]	(None, 7, 7, 64)	18496	c
=====			
flatten_1 (Flatten) conv2d_2[0][0]	(None, 3136)	0	c
=====			
dense_1 (Dense) flatten_1[0][0]	(None, 16)	50192	f
=====			
z_mean (Dense) dense_1[0][0]	(None, 2)	34	d
=====			
z_log_var (Dense) dense_1[0][0]	(None, 2)	34	d
=====			
z (Lambda) _mean[0][0]	(None, 2)	0	z
=====			
_log_var[0][0]			z
=====			
Total params: 69,076			
Trainable params: 69,076			
Non-trainable params: 0			
=====			

## Decoder with CNN

```
In [0]: latent_inputs = Input(shape=(latent_dim,), name='z_sampling')
x = Dense(shape[1] * shape[2] * shape[3], activation='relu')(latent_inputs)
x = Reshape((shape[1], shape[2], shape[3]))(x)

# use Conv2DTranspose to reverse the conv layers from the encoder
for i in range(2):
    x = Conv2DTranspose(filters=filters,
                        kernel_size=kernel_size,
                        activation='relu',
                        strides=2,
                        padding='same')(x)

    filters //= 2

outputs = Conv2DTranspose(filters=1,
                          kernel_size=kernel_size,
                          activation='sigmoid',
                          padding='same',
                          name='decoder_output')(x)

# instantiate decoder model
decoder = Model(latent_inputs, outputs, name='decoder')
decoder.summary()
```

Model: "decoder"

Layer (type)	Output Shape	Param #
z_sampling (InputLayer)	(None, 2)	0
dense_2 (Dense)	(None, 3136)	9408
reshape_1 (Reshape)	(None, 7, 7, 64)	0
conv2d_transpose_1 (Conv2DTr	(None, 14, 14, 64)	36928
conv2d_transpose_2 (Conv2DTr	(None, 28, 28, 32)	18464
decoder_output (Conv2DTransp	(None, 28, 28, 1)	289
Total params: 65,089		
Trainable params: 65,089		
Non-trainable params: 0		

## VAE

```
In [0]: outputs = decoder(encoder(inputs)[2])
vae = Model(inputs, outputs, name='vae')
```

```
In [0]: if use_mse:
        reconstruction_loss = mse(K.flatten(inputs), K.flatten(outputs)
    )
    else:
        reconstruction_loss = binary_crossentropy(K.flatten(inputs),
                                                    K.flatten(outputs))

    reconstruction_loss *= image_size * image_size
    kl_loss = 1 + z_log_var - K.square(z_mean) - K.exp(z_log_var)
    kl_loss = K.sum(kl_loss, axis=-1)
    kl_loss *= -0.5
    vae_loss = K.mean(reconstruction_loss + kl_loss)
    vae.add_loss(vae_loss)
```

```
In [0]: vae.compile(optimizer='rmsprop')
        vae.summary()
```

Model: "vae"

Layer (type)	Output Shape	Param #
encoder_input (InputLayer)	(None, 28, 28, 1)	0
encoder (Model)	[(None, 2), (None, 2), (N 69076	
decoder (Model)	(None, 28, 28, 1)	65089
Total params: 134,165		
Trainable params: 134,165		
Non-trainable params: 0		

```
/usr/local/lib/python3.6/dist-packages/keras/engine/training_utils
.py:819: UserWarning: Output decoder missing from loss dictionary.
We assume this was done on purpose. The fit and evaluate APIs will
not be expecting any data to be passed to decoder.
    'be expecting any data to be passed to {0}.'.format(name))
```

```
In [0]: vae.fit(x_train, epochs=epochs, batch_size=batch_size, validation_data
        =(x_test, None))
```


```
Train on 10000 samples, validate on 1000 samples
Epoch 1/30
10000/10000 [=====] - 8s 827us/step - los
s: 53.3987 - val_loss: 44.6226
Epoch 2/30
10000/10000 [=====] - 2s 191us/step - los
s: 44.6895 - val_loss: 42.9474
Epoch 3/30
10000/10000 [=====] - 2s 189us/step - los
s: 43.3248 - val_loss: 43.7992
Epoch 4/30
```

```
10000/10000 [=====] - 2s 192us/step - los
s: 42.6410 - val_loss: 42.1521
Epoch 5/30
10000/10000 [=====] - 2s 191us/step - los
s: 42.1135 - val_loss: 42.8118
Epoch 6/30
10000/10000 [=====] - 2s 193us/step - los
s: 41.5962 - val_loss: 41.5368
Epoch 7/30
10000/10000 [=====] - 2s 193us/step - los
s: 41.1892 - val_loss: 40.6660
Epoch 8/30
10000/10000 [=====] - 2s 190us/step - los
s: 40.7682 - val_loss: 40.8434
Epoch 9/30
10000/10000 [=====] - 2s 189us/step - los
s: 40.4491 - val_loss: 41.0781
Epoch 10/30
10000/10000 [=====] - 2s 189us/step - los
s: 40.1137 - val_loss: 40.3688
Epoch 11/30
10000/10000 [=====] - 2s 188us/step - los
s: 39.8048 - val_loss: 39.9240
Epoch 12/30
10000/10000 [=====] - 2s 190us/step - los
s: 39.5393 - val_loss: 39.9892
Epoch 13/30
10000/10000 [=====] - 2s 189us/step - los
s: 39.3133 - val_loss: 39.5158
Epoch 14/30
10000/10000 [=====] - 2s 188us/step - los
s: 39.1134 - val_loss: 40.5779
Epoch 15/30
10000/10000 [=====] - 2s 190us/step - los
s: 38.8457 - val_loss: 39.6660
Epoch 16/30
10000/10000 [=====] - 2s 187us/step - los
s: 38.6564 - val_loss: 39.4518
Epoch 17/30
10000/10000 [=====] - 2s 191us/step - los
s: 38.5165 - val_loss: 40.3426
Epoch 18/30
10000/10000 [=====] - 2s 188us/step - los
s: 38.3603 - val_loss: 39.3279
Epoch 19/30
10000/10000 [=====] - 2s 188us/step - los
s: 38.2308 - val_loss: 39.6475
Epoch 20/30
10000/10000 [=====] - 2s 190us/step - los
s: 38.0641 - val_loss: 40.2713
Epoch 21/30
10000/10000 [=====] - 2s 190us/step - los
s: 37.9180 - val_loss: 39.0292
```

```
Epoch 22/30
10000/10000 [=====] - 2s 189us/step - los
s: 37.7748 - val_loss: 38.8527
Epoch 23/30
10000/10000 [=====] - 2s 190us/step - los
s: 37.7265 - val_loss: 39.0614
Epoch 24/30
10000/10000 [=====] - 2s 193us/step - los
s: 37.6059 - val_loss: 39.2262
Epoch 25/30
10000/10000 [=====] - 2s 189us/step - los
s: 37.5234 - val_loss: 39.3531
Epoch 26/30
10000/10000 [=====] - 2s 189us/step - los
s: 37.4753 - val_loss: 39.8027
Epoch 27/30
10000/10000 [=====] - 2s 189us/step - los
s: 37.2946 - val_loss: 39.0632
Epoch 28/30
10000/10000 [=====] - 2s 196us/step - los
s: 37.2517 - val_loss: 39.3951
Epoch 29/30
10000/10000 [=====] - 2s 196us/step - los
s: 37.1672 - val_loss: 39.3845
Epoch 30/30
10000/10000 [=====] - 2s 198us/step - los
s: 37.0517 - val_loss: 40.3112
```

Out[0]: <keras.callbacks.callbacks.History at 0x7f12102de438>

```
In [0]: plot_results(encoder, decoder, x_test, y_test, batch_size=batch_siz
e)
```



## Problem 4

Generative Adversarial Network consists of two models: a **generative model** G and a **discriminative model** D. The discriminator is responsible for distinguishing between the actual and generated data while the generator tries creating data in a way to fool the discriminator.

On using cross entropy loss function:

$$L(y, \hat{y}) = (y \log \hat{y}) + ((1 - y) \log(1 - \hat{y}))$$

While training the discriminator label of data coming from  $P_{data}(x)$  is  $y=1$  and  $\hat{y} = D(x)$  thus the loss we get is:

$$L(D(x), 1) = \log(D(x))$$

Since the data for the generator is fake, therefore  $y=0$  and  $\hat{y} = D(G(z))$

$$\text{So } L(D(G(z)), 0) = \log(1 - D(G(z)))$$

The Objective of discriminator is to classify real and fake data.

The loss function of the generator model is :-

$$L^{(G)} = \min[\log(D(x)) + \log(1 - D(G(z)))]$$

The loss function of the Discriminator model is :-

$$L^{(D)} = \max[\log(D(x)) + \log(1 - D(G(z)))]$$

Therefore the **Combined loss function** is:

$$L = \min_G \max_D [E_{x \sim P_{data}(x)} [\log(D(x))] + E_{z \sim P(z)} [\log(1 - D(G(z)))]]$$

Thus, we can update the discriminator model based on ascending the stochastic gradient descent while update the generator model by descending the stochastic gradient descent.



```
In [0]: import numpy as np
import torch
import matplotlib.pyplot as plt
from torchvision import datasets
import torchvision.transforms as transforms

# number of subprocesses to use for data loading
num_workers = 0
# how many samples per batch to load
batch_size = 64

# convert data to torch.FloatTensor
transform = transforms.ToTensor()

# get the training datasets
train_data = datasets.MNIST(root='data', train=True,
                             download=True, transform=transfo
rm)

# prepare data loader
train_loader = torch.utils.data.DataLoader(train_data, batch_size=b
atch_size,
                                             num_workers=num_workers)
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to data/MNIST/raw/train-images-idx3-ubyte.gz

Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to data/MNIST/raw/train-labels-idx1-ubyte.gz

Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to data/MNIST/raw/t10k-images-idx3-ubyte.gz

Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to data/MNIST/raw/t10k-labels-idx1-ubyte.gz

Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw

Processing...

Done!

## Discriminator

```
In [0]: import torch.nn as nn
import torch.nn.functional as F

class Discriminator(nn.Module):

    def __init__(self, input_size, hidden_dim, output_size):
        super(Discriminator, self).__init__()

        # define hidden linear layers
        self.fc1 = nn.Linear(input_size, hidden_dim*4)
        self.fc2 = nn.Linear(hidden_dim*4, hidden_dim*2)
        self.fc3 = nn.Linear(hidden_dim*2, hidden_dim)

        # final fully-connected layer
        self.fc4 = nn.Linear(hidden_dim, output_size)

        # dropout layer
        self.dropout = nn.Dropout(0.3)

    def forward(self, x):
        # flatten image
        x = x.view(-1, 28*28)
        # all hidden layers
        x = F.leaky_relu(self.fc1(x), 0.2) # (input, negative_slope
=0.2)
        x = self.dropout(x)
        x = F.leaky_relu(self.fc2(x), 0.2)
        x = self.dropout(x)
        x = F.leaky_relu(self.fc3(x), 0.2)
        x = self.dropout(x)
        # final layer
        out = self.fc4(x)

        return out
```

## Generator

```
In [0]: class Generator(nn.Module):

    def __init__(self, input_size, hidden_dim, output_size):
        super(Generator, self).__init__()

        # define hidden linear layers
        self.fc1 = nn.Linear(input_size, hidden_dim)
        self.fc2 = nn.Linear(hidden_dim, hidden_dim*2)
        self.fc3 = nn.Linear(hidden_dim*2, hidden_dim*4)

        # final fully-connected layer
        self.fc4 = nn.Linear(hidden_dim*4, output_size)

        # dropout layer
        self.dropout = nn.Dropout(0.3)

    def forward(self, x):
        # all hidden layers
        x = F.leaky_relu(self.fc1(x), 0.2) # (input, negative_slope
=0.2)
        x = self.dropout(x)
        x = F.leaky_relu(self.fc2(x), 0.2)
        x = self.dropout(x)
        x = F.leaky_relu(self.fc3(x), 0.2)
        x = self.dropout(x)
        # final layer with tanh applied
        out = torch.tanh(self.fc4(x))

        return out
```

## Hyperparameters

```
In [0]: # Discriminator hyperparams

# Size of input image to discriminator (28*28)
input_size = 784
# Size of discriminator output (real or fake)
d_output_size = 1
# Size of last hidden layer in the discriminator
d_hidden_size = 32

# Generator hyperparams

# Size of latent vector to give to generator
z_size = 100
# Size of discriminator output (generated image)
g_output_size = 784
# Size of first hidden layer in the generator
g_hidden_size = 32
```

## GAN model

```
In [0]: D = Discriminator(input_size, d_hidden_size, d_output_size)
        G = Generator(z_size, g_hidden_size, g_output_size)

        print(D)
        print()
        print(G)

        Discriminator(
            (fc1): Linear(in_features=784, out_features=128, bias=True)
            (fc2): Linear(in_features=128, out_features=64, bias=True)
            (fc3): Linear(in_features=64, out_features=32, bias=True)
            (fc4): Linear(in_features=32, out_features=1, bias=True)
            (dropout): Dropout(p=0.3, inplace=False)
        )

        Generator(
            (fc1): Linear(in_features=100, out_features=32, bias=True)
            (fc2): Linear(in_features=32, out_features=64, bias=True)
            (fc3): Linear(in_features=64, out_features=128, bias=True)
            (fc4): Linear(in_features=128, out_features=784, bias=True)
            (dropout): Dropout(p=0.3, inplace=False)
        )
```

```
In [0]: # Calculate losses
        def real_loss(D_out, smooth=False):
            batch_size = D_out.size(0)
            # label smoothing
            if smooth:
                # smooth, real labels = 0.9
                labels = torch.ones(batch_size)*0.9
            else:
                labels = torch.ones(batch_size) # real labels = 1

            # numerically stable loss
            criterion = nn.BCEWithLogitsLoss()
            # calculate loss
            loss = criterion(D_out.squeeze(), labels)
            return loss

        def fake_loss(D_out):
            batch_size = D_out.size(0)
            labels = torch.zeros(batch_size) # fake labels = 0
            criterion = nn.BCEWithLogitsLoss()
            # calculate loss
            loss = criterion(D_out.squeeze(), labels)
            return loss
```

```
In [0]: import torch.optim as optim

# Optimizers
lr = 0.002

# Create optimizers for the discriminator and generator
d_optimizer = optim.Adam(D.parameters(), lr)
g_optimizer = optim.Adam(G.parameters(), lr)
```

## Training

```
In [0]: # training hyperparams
num_epochs = 50

# keep track of loss and generated, "fake" samples
samples = []
losses = []

print_every = 400

# Get some fixed data for sampling. These are images that are held
# constant throughout training, and allow us to inspect the model's
# performance
sample_size=16
fixed_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
fixed_z = torch.from_numpy(fixed_z).float()

# train the network
D.train()
G.train()
for epoch in range(num_epochs):

    for batch_i, (real_images, _) in enumerate(train_loader):

        batch_size = real_images.size(0)

        ## Important rescaling step ##
        real_images = real_images*2 - 1 # rescale input images from
        [0,1) to [-1, 1)

        d_optimizer.zero_grad()

        # 1. Train with real images

        # Compute the discriminator losses on real images
        # smooth the real labels
        D_real = D(real_images)
        d_real_loss = real_loss(D_real, smooth=True)

        # 2. Train with fake images
```

```

# Generate fake images
z = np.random.uniform(-1, 1, size=(batch_size, z_size))
z = torch.from_numpy(z).float()
fake_images = G(z)

# Compute the discriminator losses on fake images
D_fake = D(fake_images)
d_fake_loss = fake_loss(D_fake)

# add up loss and perform backprop
d_loss = d_real_loss + d_fake_loss
d_loss.backward()
d_optimizer.step()

g_optimizer.zero_grad()

# 1. Train with fake images and flipped labels

# Generate fake images
z = np.random.uniform(-1, 1, size=(batch_size, z_size))
z = torch.from_numpy(z).float()
fake_images = G(z)

# Compute the discriminator losses on fake images
# using flipped labels!
D_fake = D(fake_images)
g_loss = real_loss(D_fake) # use real loss to flip labels

# perform backprop
g_loss.backward()
g_optimizer.step()

# Print some loss stats
if batch_i % print_every == 0:
    # print discriminator and generator loss
    print('Epoch [{:5d}/{:5d}] | d_loss: {:.64f} | g_loss:
{:6.4f}'.format(
        epoch+1, num_epochs, d_loss.item(), g_loss.item
    )))

## AFTER EACH EPOCH##
# append discriminator loss and generator loss
losses.append((d_loss.item(), g_loss.item()))

# generate and save sample, fake images
G.eval() # eval mode for generating samples
samples_z = G(fixed_z)
samples.append(samples_z)
G.train() # back to train mode

```

```

Epoch [    1/   50] | d_loss: 1.4199 | g_loss: 0.6122
Epoch [    1/   50] | d_loss: 1.6515 | g_loss: 1.4794

```

Epoch [	1/	50]		d_loss: 1.0764		g_loss: 2.1332
Epoch [	2/	50]		d_loss: 1.1537		g_loss: 1.1083
Epoch [	2/	50]		d_loss: 1.2909		g_loss: 0.6969
Epoch [	2/	50]		d_loss: 1.6127		g_loss: 0.5472
Epoch [	3/	50]		d_loss: 1.2701		g_loss: 0.8140
Epoch [	3/	50]		d_loss: 1.1618		g_loss: 0.9189
Epoch [	3/	50]		d_loss: 1.2258		g_loss: 2.5967
Epoch [	4/	50]		d_loss: 1.2586		g_loss: 1.2860
Epoch [	4/	50]		d_loss: 1.0651		g_loss: 1.5339
Epoch [	4/	50]		d_loss: 1.0524		g_loss: 1.2302
Epoch [	5/	50]		d_loss: 1.1280		g_loss: 1.0989
Epoch [	5/	50]		d_loss: 1.2278		g_loss: 1.2629
Epoch [	5/	50]		d_loss: 1.4973		g_loss: 0.7280
Epoch [	6/	50]		d_loss: 1.2898		g_loss: 1.1424
Epoch [	6/	50]		d_loss: 1.3041		g_loss: 0.9609
Epoch [	6/	50]		d_loss: 1.4286		g_loss: 0.8839
Epoch [	7/	50]		d_loss: 1.1804		g_loss: 1.1715
Epoch [	7/	50]		d_loss: 1.1365		g_loss: 1.2973
Epoch [	7/	50]		d_loss: 1.0273		g_loss: 2.4373
Epoch [	8/	50]		d_loss: 1.2099		g_loss: 1.6809
Epoch [	8/	50]		d_loss: 1.3877		g_loss: 1.2304
Epoch [	8/	50]		d_loss: 1.2133		g_loss: 2.4702
Epoch [	9/	50]		d_loss: 1.1091		g_loss: 1.0650
Epoch [	9/	50]		d_loss: 1.2882		g_loss: 1.0777
Epoch [	9/	50]		d_loss: 1.2394		g_loss: 1.6353
Epoch [	10/	50]		d_loss: 1.3832		g_loss: 1.0669
Epoch [	10/	50]		d_loss: 1.2868		g_loss: 1.0736
Epoch [	10/	50]		d_loss: 1.3039		g_loss: 0.7582
Epoch [	11/	50]		d_loss: 1.2426		g_loss: 1.6015
Epoch [	11/	50]		d_loss: 1.0400		g_loss: 1.0918
Epoch [	11/	50]		d_loss: 1.2344		g_loss: 1.1228
Epoch [	12/	50]		d_loss: 1.1862		g_loss: 1.5529
Epoch [	12/	50]		d_loss: 1.0059		g_loss: 1.6859
Epoch [	12/	50]		d_loss: 1.1994		g_loss: 1.0489
Epoch [	13/	50]		d_loss: 1.2884		g_loss: 0.8636
Epoch [	13/	50]		d_loss: 1.2170		g_loss: 0.8706
Epoch [	13/	50]		d_loss: 1.3117		g_loss: 1.1369
Epoch [	14/	50]		d_loss: 1.2062		g_loss: 1.4767
Epoch [	14/	50]		d_loss: 1.3353		g_loss: 1.1355
Epoch [	14/	50]		d_loss: 1.3380		g_loss: 0.8093
Epoch [	15/	50]		d_loss: 1.3188		g_loss: 1.0626
Epoch [	15/	50]		d_loss: 1.2285		g_loss: 1.0403
Epoch [	15/	50]		d_loss: 1.2061		g_loss: 1.3532
Epoch [	16/	50]		d_loss: 1.1674		g_loss: 1.1758
Epoch [	16/	50]		d_loss: 1.2690		g_loss: 1.5118
Epoch [	16/	50]		d_loss: 1.1016		g_loss: 1.4641
Epoch [	17/	50]		d_loss: 1.1594		g_loss: 1.3556
Epoch [	17/	50]		d_loss: 1.3024		g_loss: 1.1883
Epoch [	17/	50]		d_loss: 1.3620		g_loss: 0.9748
Epoch [	18/	50]		d_loss: 1.2076		g_loss: 1.1343
Epoch [	18/	50]		d_loss: 1.3130		g_loss: 0.9361
Epoch [	18/	50]		d_loss: 1.2873		g_loss: 1.1706
Epoch [	19/	50]		d_loss: 1.5852		g_loss: 1.2933

Epoch [	19/	50]	d_loss: 1.1843	g_loss: 1.0082
Epoch [	19/	50]	d_loss: 1.3097	g_loss: 1.1157
Epoch [	20/	50]	d_loss: 1.2634	g_loss: 0.9743
Epoch [	20/	50]	d_loss: 1.1606	g_loss: 1.1301
Epoch [	20/	50]	d_loss: 1.3761	g_loss: 1.1775
Epoch [	21/	50]	d_loss: 1.2424	g_loss: 0.9947
Epoch [	21/	50]	d_loss: 1.1746	g_loss: 1.0627
Epoch [	21/	50]	d_loss: 1.3331	g_loss: 0.9228
Epoch [	22/	50]	d_loss: 1.2118	g_loss: 1.1222
Epoch [	22/	50]	d_loss: 1.1306	g_loss: 1.6337
Epoch [	22/	50]	d_loss: 1.2765	g_loss: 1.1879
Epoch [	23/	50]	d_loss: 1.2515	g_loss: 0.9422
Epoch [	23/	50]	d_loss: 1.1771	g_loss: 1.3238
Epoch [	23/	50]	d_loss: 1.4102	g_loss: 1.1210
Epoch [	24/	50]	d_loss: 1.2850	g_loss: 1.0056
Epoch [	24/	50]	d_loss: 1.2490	g_loss: 1.0430
Epoch [	24/	50]	d_loss: 1.3396	g_loss: 0.8730
Epoch [	25/	50]	d_loss: 1.2385	g_loss: 1.2466
Epoch [	25/	50]	d_loss: 1.2371	g_loss: 1.5061
Epoch [	25/	50]	d_loss: 1.3642	g_loss: 0.8888
Epoch [	26/	50]	d_loss: 1.3380	g_loss: 1.1839
Epoch [	26/	50]	d_loss: 1.1741	g_loss: 1.0695
Epoch [	26/	50]	d_loss: 1.3976	g_loss: 1.0299
Epoch [	27/	50]	d_loss: 1.3432	g_loss: 0.8521
Epoch [	27/	50]	d_loss: 1.1745	g_loss: 1.0786
Epoch [	27/	50]	d_loss: 1.3620	g_loss: 1.0966
Epoch [	28/	50]	d_loss: 1.2902	g_loss: 1.0683
Epoch [	28/	50]	d_loss: 1.2502	g_loss: 1.3421
Epoch [	28/	50]	d_loss: 1.2578	g_loss: 1.2000
Epoch [	29/	50]	d_loss: 1.3181	g_loss: 1.1824
Epoch [	29/	50]	d_loss: 1.2371	g_loss: 0.8801
Epoch [	29/	50]	d_loss: 1.4353	g_loss: 0.9504
Epoch [	30/	50]	d_loss: 1.3651	g_loss: 1.2938
Epoch [	30/	50]	d_loss: 1.1633	g_loss: 1.0334
Epoch [	30/	50]	d_loss: 1.3691	g_loss: 1.0443
Epoch [	31/	50]	d_loss: 1.2872	g_loss: 1.0678
Epoch [	31/	50]	d_loss: 1.3934	g_loss: 1.0067
Epoch [	31/	50]	d_loss: 1.3867	g_loss: 1.0341
Epoch [	32/	50]	d_loss: 1.2284	g_loss: 1.1083
Epoch [	32/	50]	d_loss: 1.1970	g_loss: 0.9522
Epoch [	32/	50]	d_loss: 1.3643	g_loss: 0.9842
Epoch [	33/	50]	d_loss: 1.3168	g_loss: 0.9101
Epoch [	33/	50]	d_loss: 1.2395	g_loss: 0.9652
Epoch [	33/	50]	d_loss: 1.3533	g_loss: 1.0110
Epoch [	34/	50]	d_loss: 1.2560	g_loss: 0.9041
Epoch [	34/	50]	d_loss: 1.2547	g_loss: 1.1997
Epoch [	34/	50]	d_loss: 1.3714	g_loss: 1.0845
Epoch [	35/	50]	d_loss: 1.2783	g_loss: 1.1826
Epoch [	35/	50]	d_loss: 1.2055	g_loss: 1.1693
Epoch [	35/	50]	d_loss: 1.4162	g_loss: 0.9825
Epoch [	36/	50]	d_loss: 1.2335	g_loss: 1.0916
Epoch [	36/	50]	d_loss: 1.3161	g_loss: 1.1513
Epoch [	36/	50]	d_loss: 1.3667	g_loss: 0.9204



Epoch	[	37/	50]		d_loss:	1.3909		g_loss:	0.9699
Epoch	[	37/	50]		d_loss:	1.1925		g_loss:	0.8976
Epoch	[	37/	50]		d_loss:	1.2214		g_loss:	1.1316
Epoch	[	38/	50]		d_loss:	1.2759		g_loss:	1.4615
Epoch	[	38/	50]		d_loss:	1.2687		g_loss:	1.1092
Epoch	[	38/	50]		d_loss:	1.3101		g_loss:	1.0101
Epoch	[	39/	50]		d_loss:	1.2209		g_loss:	1.0353
Epoch	[	39/	50]		d_loss:	1.3207		g_loss:	0.8350
Epoch	[	39/	50]		d_loss:	1.3460		g_loss:	1.0167
Epoch	[	40/	50]		d_loss:	1.2809		g_loss:	1.1309
Epoch	[	40/	50]		d_loss:	1.1904		g_loss:	1.3867
Epoch	[	40/	50]		d_loss:	1.2163		g_loss:	0.9536
Epoch	[	41/	50]		d_loss:	1.2880		g_loss:	1.2844
Epoch	[	41/	50]		d_loss:	1.1887		g_loss:	1.1490
Epoch	[	41/	50]		d_loss:	1.3812		g_loss:	0.9678
Epoch	[	42/	50]		d_loss:	1.2835		g_loss:	1.0684
Epoch	[	42/	50]		d_loss:	1.3359		g_loss:	0.9492
Epoch	[	42/	50]		d_loss:	1.2633		g_loss:	0.8281
Epoch	[	43/	50]		d_loss:	1.2822		g_loss:	0.9608
Epoch	[	43/	50]		d_loss:	1.2924		g_loss:	1.2194
Epoch	[	43/	50]		d_loss:	1.3848		g_loss:	1.0217
Epoch	[	44/	50]		d_loss:	1.1986		g_loss:	1.0440
Epoch	[	44/	50]		d_loss:	1.2622		g_loss:	1.1229
Epoch	[	44/	50]		d_loss:	1.3113		g_loss:	0.8368
Epoch	[	45/	50]		d_loss:	1.2967		g_loss:	0.9703
Epoch	[	45/	50]		d_loss:	1.2194		g_loss:	1.1371
Epoch	[	45/	50]		d_loss:	1.3074		g_loss:	1.0420
Epoch	[	46/	50]		d_loss:	1.2673		g_loss:	1.1737
Epoch	[	46/	50]		d_loss:	1.1960		g_loss:	1.2485
Epoch	[	46/	50]		d_loss:	1.3350		g_loss:	1.2967
Epoch	[	47/	50]		d_loss:	1.3494		g_loss:	1.0837
Epoch	[	47/	50]		d_loss:	1.2017		g_loss:	1.0485
Epoch	[	47/	50]		d_loss:	1.3004		g_loss:	0.9711
Epoch	[	48/	50]		d_loss:	1.2758		g_loss:	0.9524
Epoch	[	48/	50]		d_loss:	1.2735		g_loss:	1.1616
Epoch	[	48/	50]		d_loss:	1.3285		g_loss:	1.1261
Epoch	[	49/	50]		d_loss:	1.1585		g_loss:	1.2917
Epoch	[	49/	50]		d_loss:	1.2306		g_loss:	1.2467
Epoch	[	49/	50]		d_loss:	1.2748		g_loss:	1.0468
Epoch	[	50/	50]		d_loss:	1.2563		g_loss:	0.9532
Epoch	[	50/	50]		d_loss:	1.2174		g_loss:	0.9255
Epoch	[	50/	50]		d_loss:	1.2864		g_loss:	0.8795

```
In [0]: losses = np.array(losses)
plt.plot(losses.T[0], label='Discriminator')
plt.plot(losses.T[1], label='Generator')
plt.title("Training Losses")
plt.legend()
```

Out[0]: <matplotlib.legend.Legend at 0x7f224e8afbe0>



```
In [0]: def view_samples(epoch, samples):
fig, axes = plt.subplots(figsize=(7,7), nrows=4, ncols=4, share
y=True, sharex=True)
for ax, img in zip(axes.flatten(), samples[epoch]):
img = img.detach()
ax.xaxis.set_visible(False)
ax.yaxis.set_visible(False)
im = ax.imshow(img.reshape((28,28)), cmap='Greys_r')
```

```
In [0]: # -1 indicates final epoch's samples (the last in the list)
view_samples(-1, samples)
```



```
In [0]: # randomly generated, new latent vectors
sample_size=16
rand_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
rand_z = torch.from_numpy(rand_z).float()

G.eval() # eval mode
# generated samples
rand_images = G(rand_z)

# 0 indicates the first set of samples in the passed in list
# and we only have one batch of samples, here
view_samples(0, [rand_images])
```



**References :**

<https://www.kaggle.com/vincentman0403/vae-with-convolution-on-mnist>  
(<https://www.kaggle.com/vincentman0403/vae-with-convolution-on-mnist>)

<https://medium.com/intel-student-ambassadors/mnist-gan-detailed-step-by-step-explanation-implementation-in-code-ecc93b22dc60> (<https://medium.com/intel-student-ambassadors/mnist-gan-detailed-step-by-step-explanation-implementation-in-code-ecc93b22dc60>)

[https://keras.io/examples/variational\\_autoencoder/](https://keras.io/examples/variational_autoencoder/) ([https://keras.io/examples/variational\\_autoencoder/](https://keras.io/examples/variational_autoencoder/))

<http://lyy1994.github.io/machine-learning/2017/04/17/RBM-tensorflow-implementation.html>  
(<http://lyy1994.github.io/machine-learning/2017/04/17/RBM-tensorflow-implementation.html>)

<https://wiseodd.github.io/techblog/2016/12/10/variational-autoencoder/>  
(<https://wiseodd.github.io/techblog/2016/12/10/variational-autoencoder/>)

<http://lyy1994.github.io/machine-learning/2017/03/16/EBM-Notes.html>  
(<http://lyy1994.github.io/machine-learning/2017/03/16/EBM-Notes.html>)